Speculative Decoding

Visual diagnosis of wall time and speed-up limits

- Data source: `results/all_configs_summary.csv` and `results/pivot_time_breakdown_debug_quickrun.csv`
- Focus: `alpha`, `B_eff`, `k`, and stage-wise bottlenecks

What We Observe First

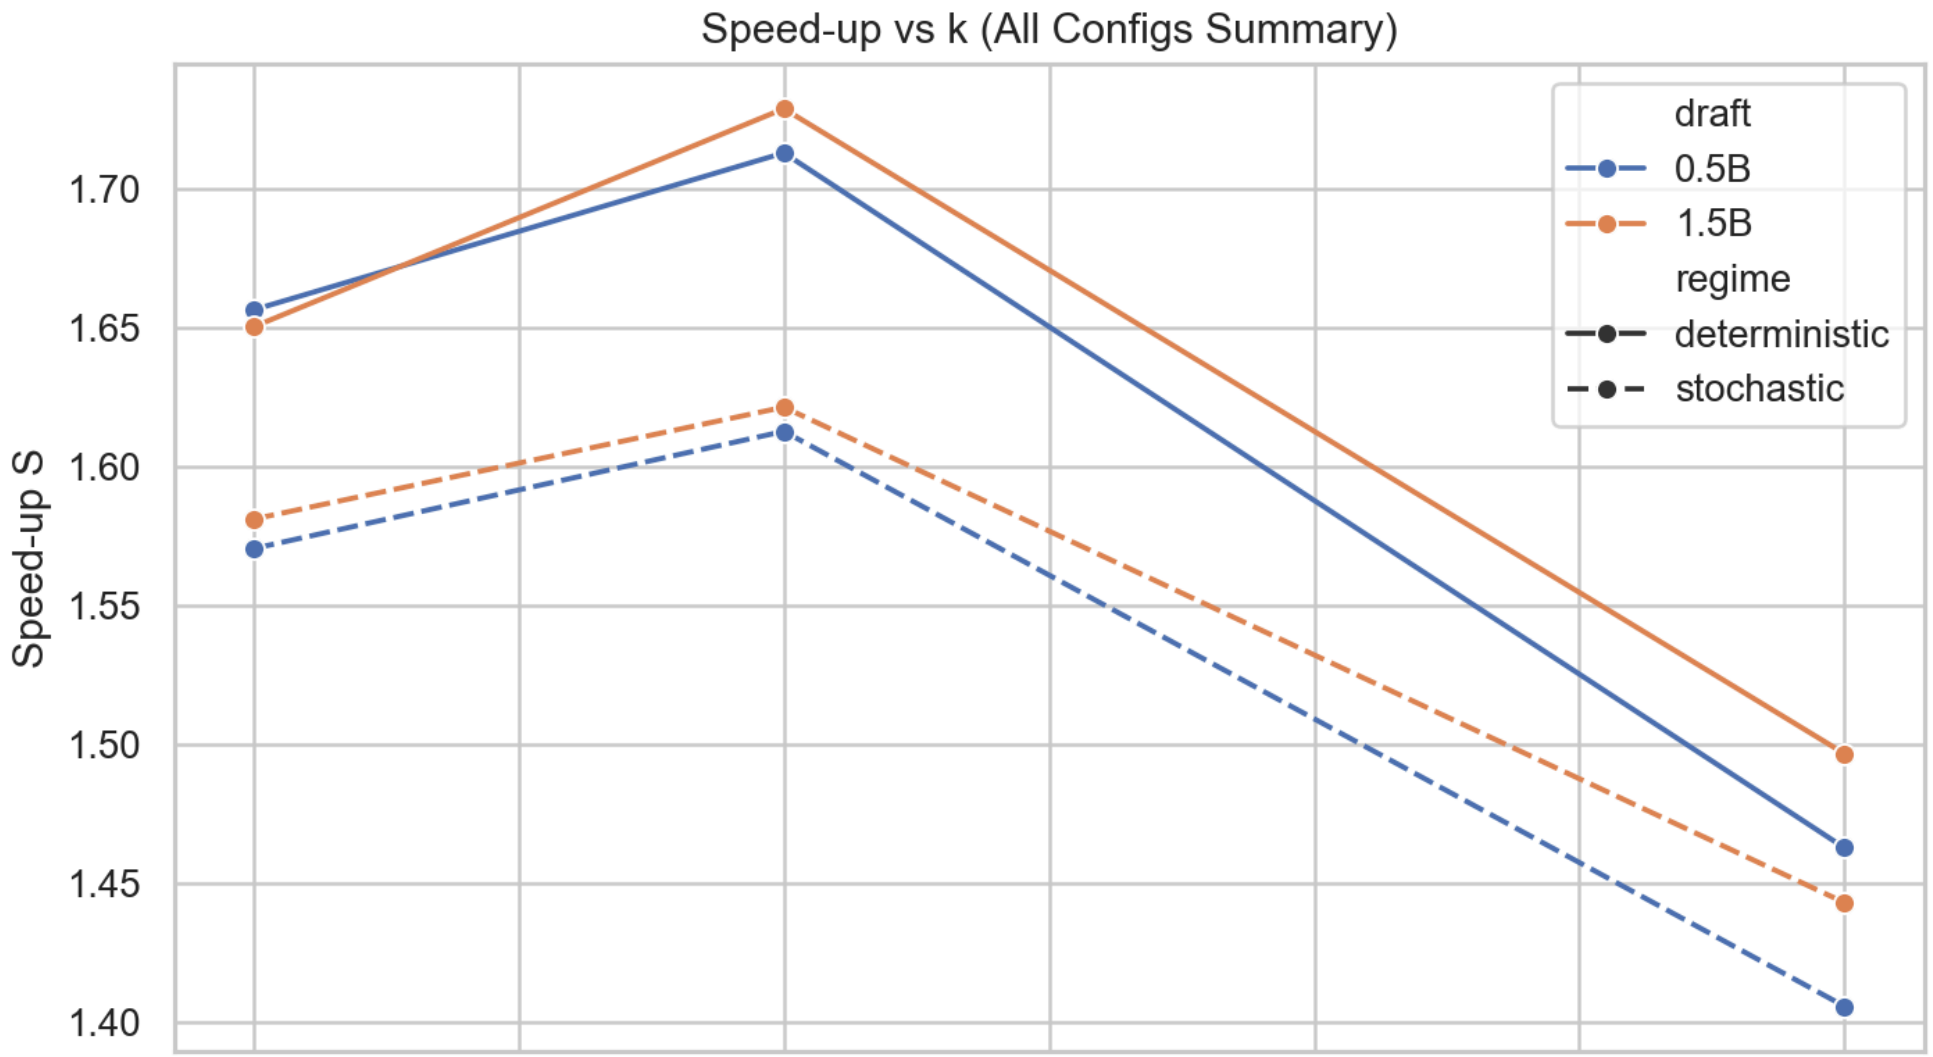

Acceptance Dynamics

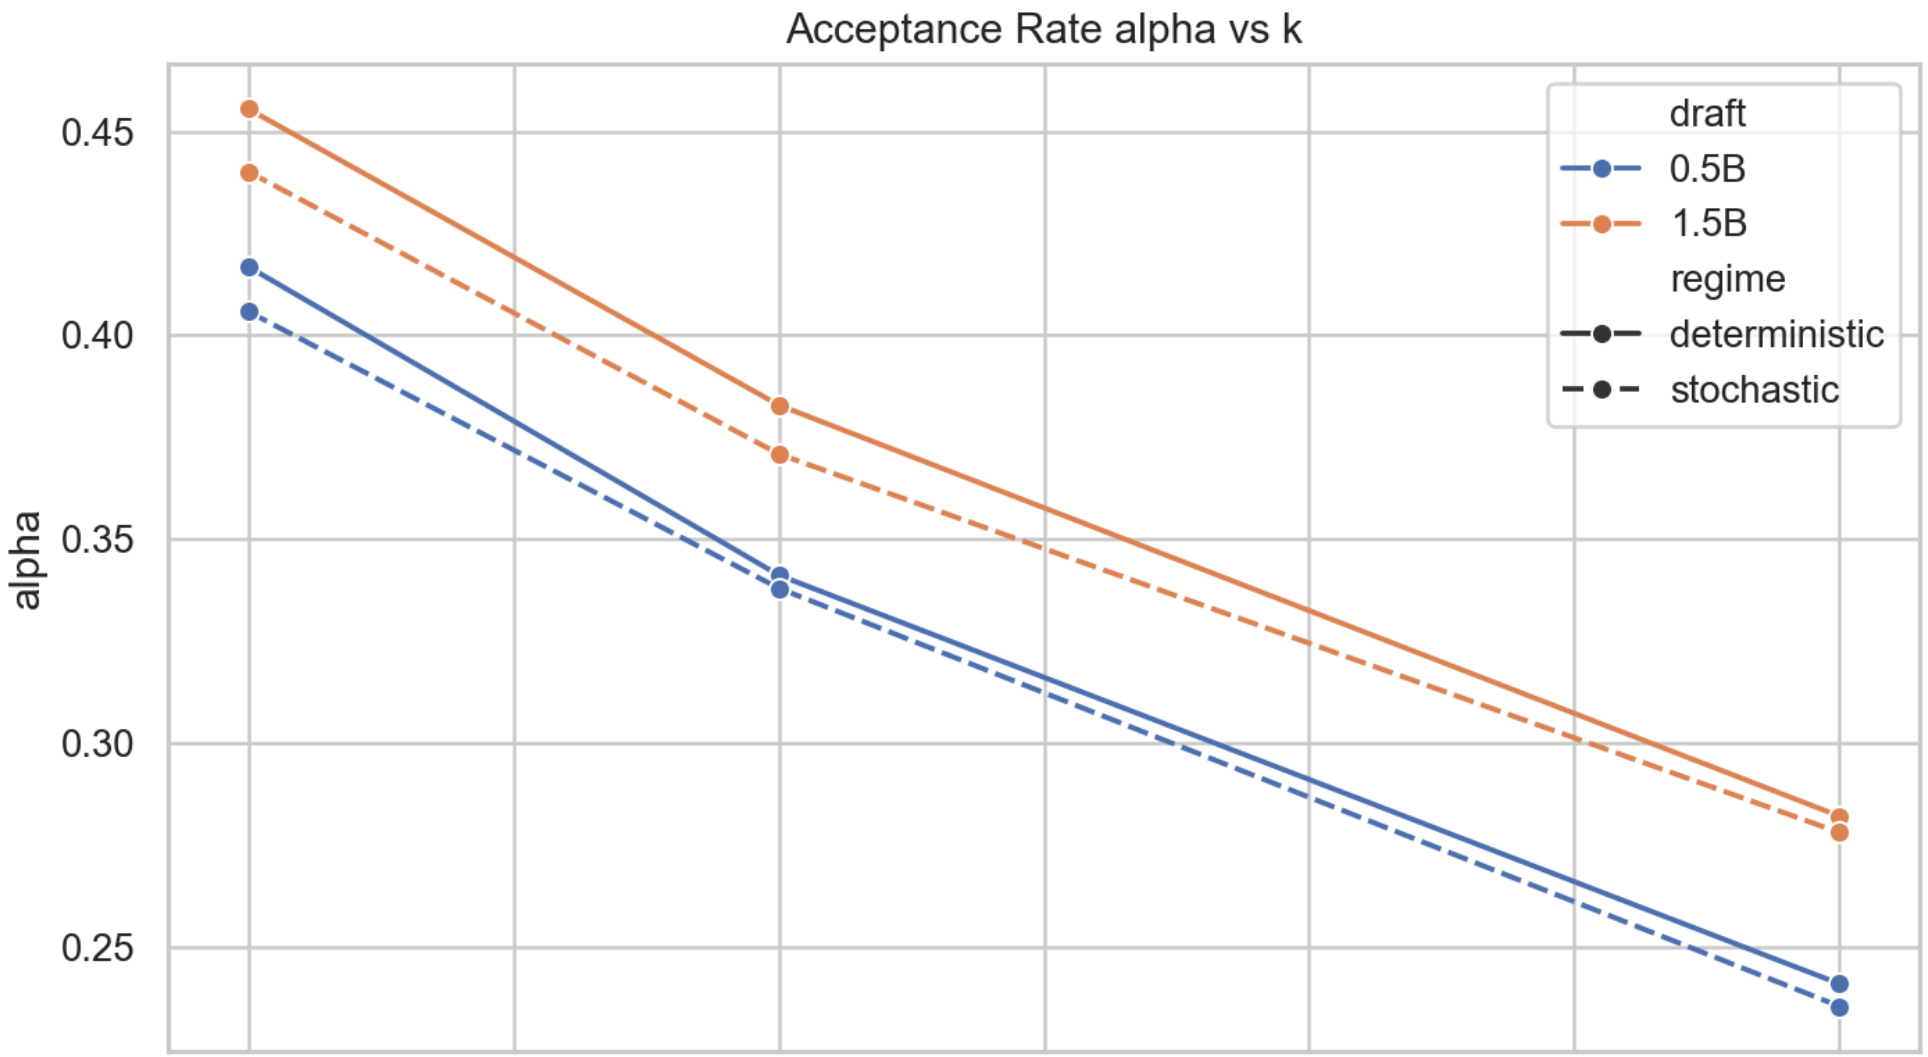

Where Time Is Spent (Share)

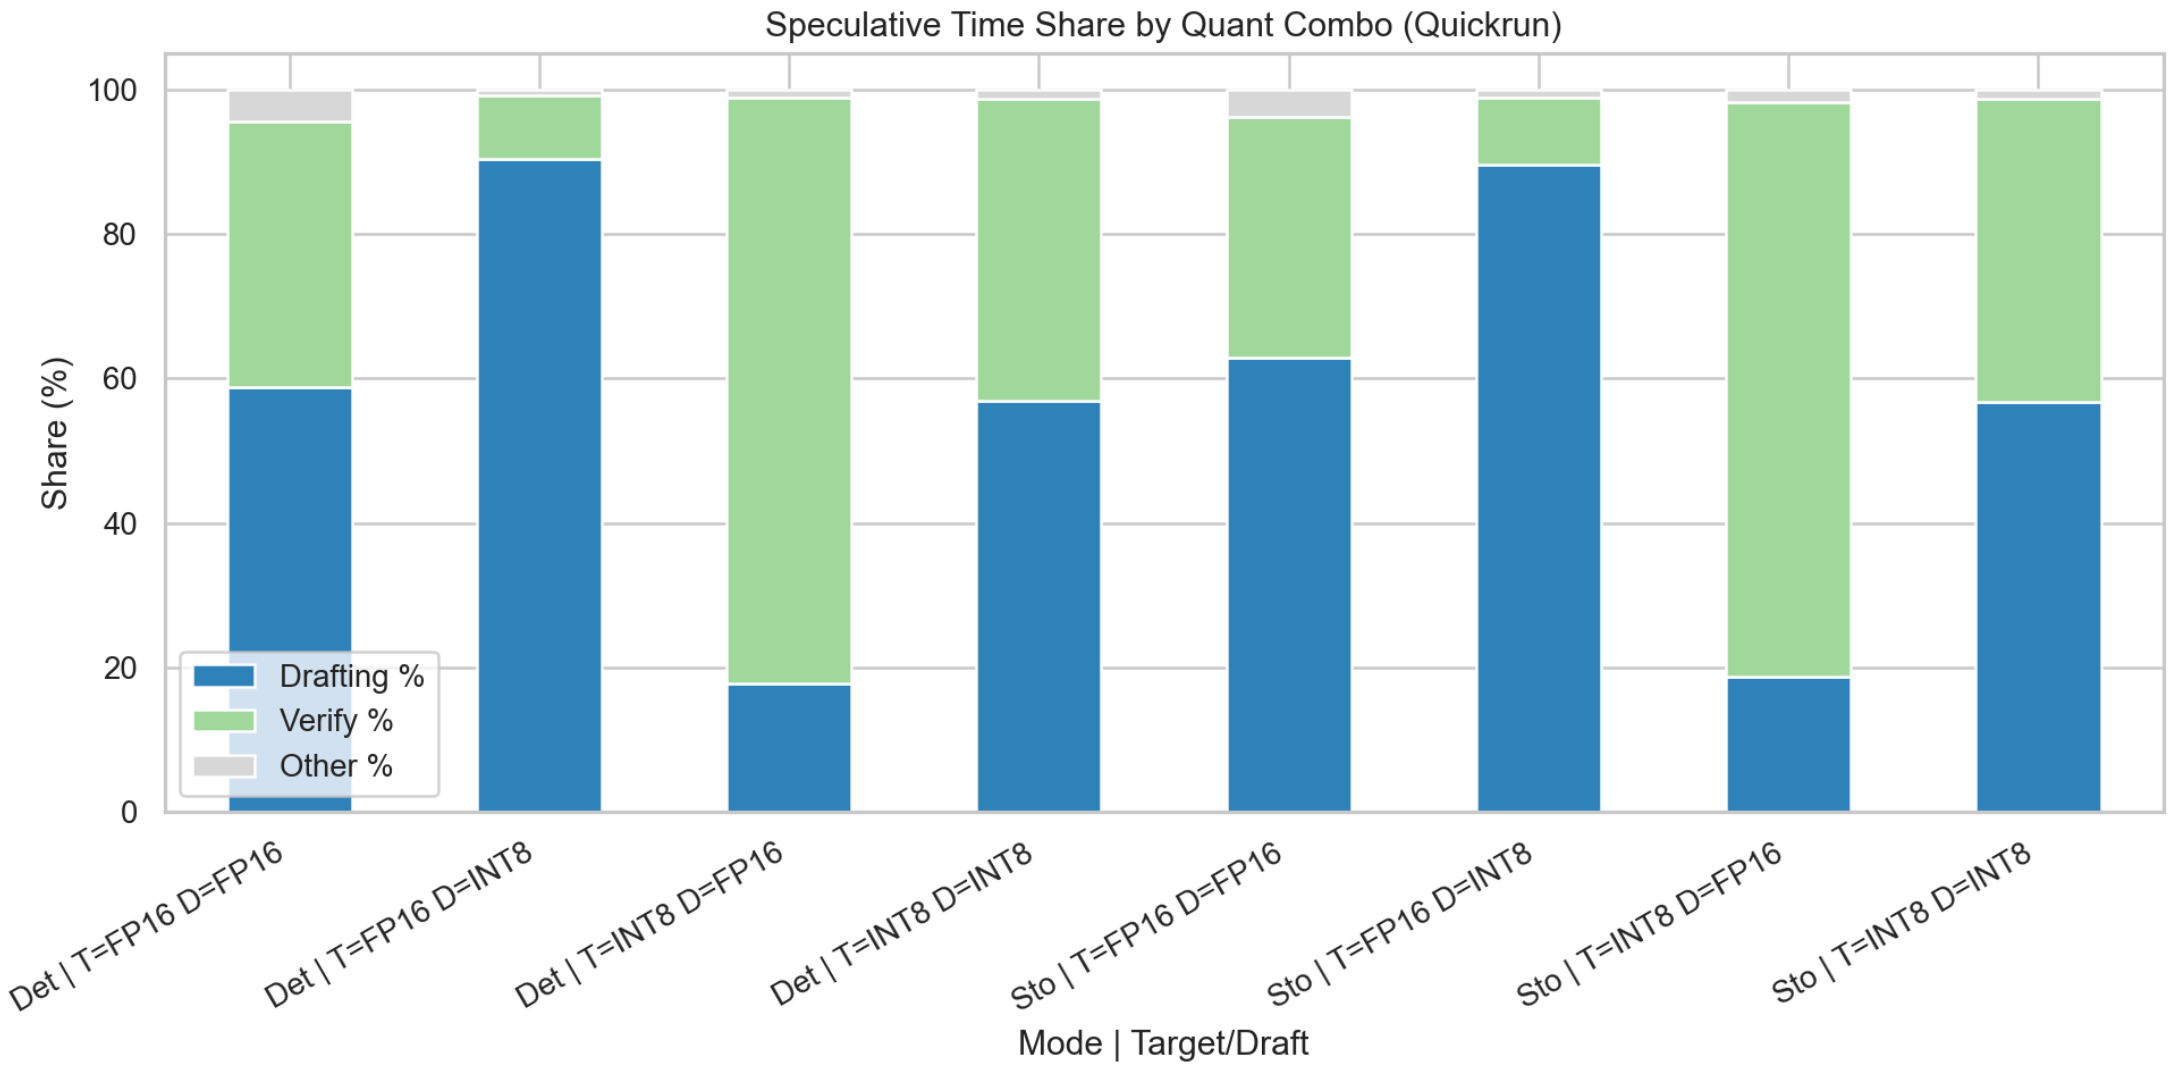

Where Time Is Spent (Seconds)

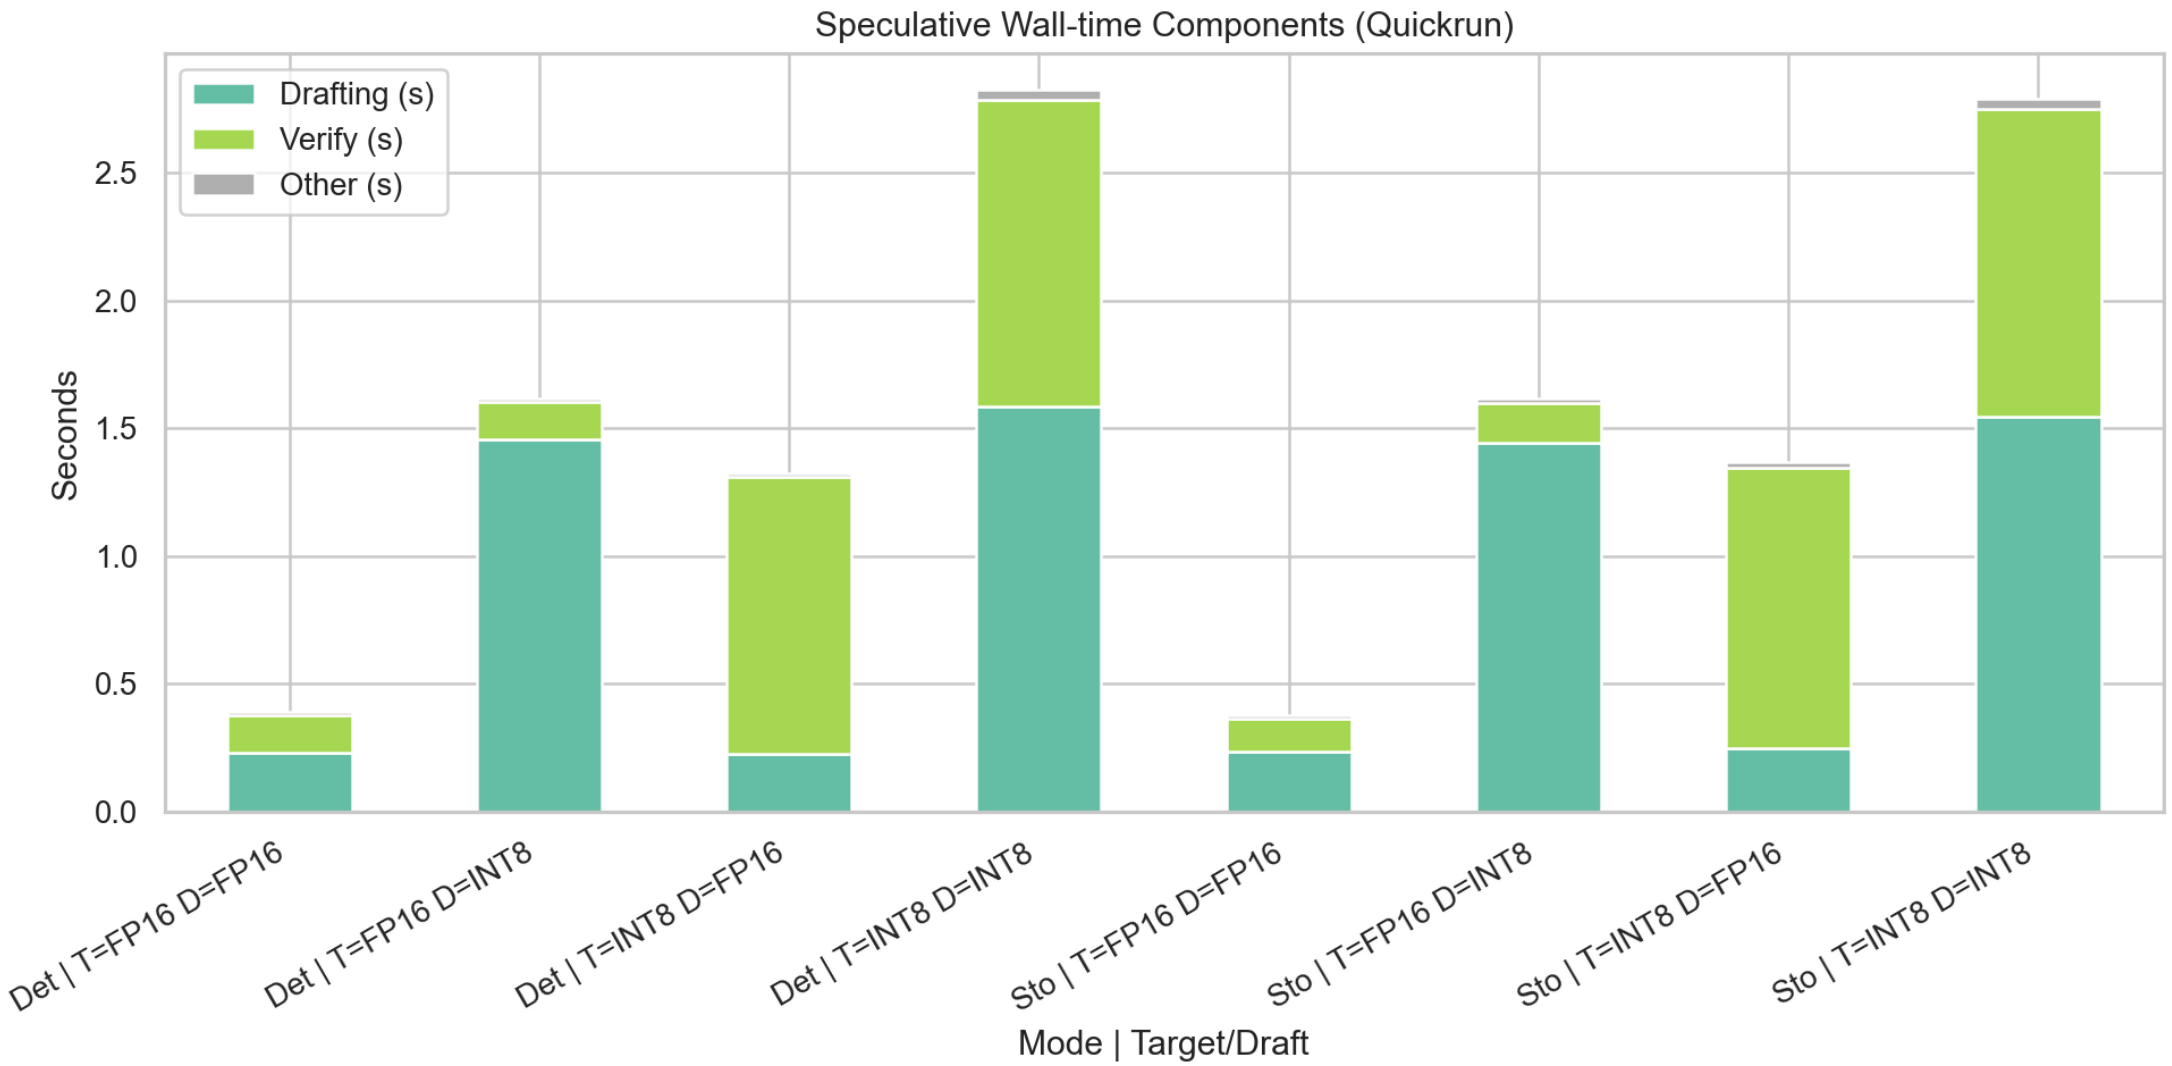

Metric Lens

$$\alpha = \frac{\text{accepted}}{\text{proposed}}, \qquad B_{\text{eff}} = \frac{\text{accepted}}{\text{verify steps}}$$

$$S = \frac{T_{\text{baseline}}}{T_{\text{spec}}}, \qquad T_{\text{spec}} \approx T_{\text{draft}} + T_{\text{verify}} + T_{\text{other}}$$

- Even good `alpha` can underperform if draft/verify cost is imbalanced.

Why FP16 + FP16 Can Still Underperform

- Moderate acceptance still causes correction overhead.
- Draft path and verify path can remain serialized on one GPU.
- Fixed costs (prefill, synchronization, kernel launch) dominate short outputs.
- `k` above optimum increases wasted draft work.

Practical Implication

- Optimize the system, not a single precision knob.
- Priorities:
  i. Find the local-optimal `k`.
  ii. Minimize dominant stage (draft or verify).
  iii. Benchmark on longer decode lengths to reduce fixed-cost distortion.
  iv. Track per-stage shares together with `alpha` and `B_eff`.

Next Step

Run compact ablation over:

- `k` (proposal length)
- output length (`max_new_tokens`)
- prompt length bucket (short/medium/long)

Goal: quantify which factor contributes most to wall time and speed-up variance.