


CUDA optimisation toggle benchmark on RTX 5090 Laptop
Output match: 100.0% (n=12 paired samples)

Overall effect with 95% bootstrap CI



Task-level latency delta

